C Resumed TLS 1.3 (four suites), predictions fixed in advance

run_testbed --resume; 222-byte response, so only the client's request flight can cost a round trip; n = 30; ~25 min

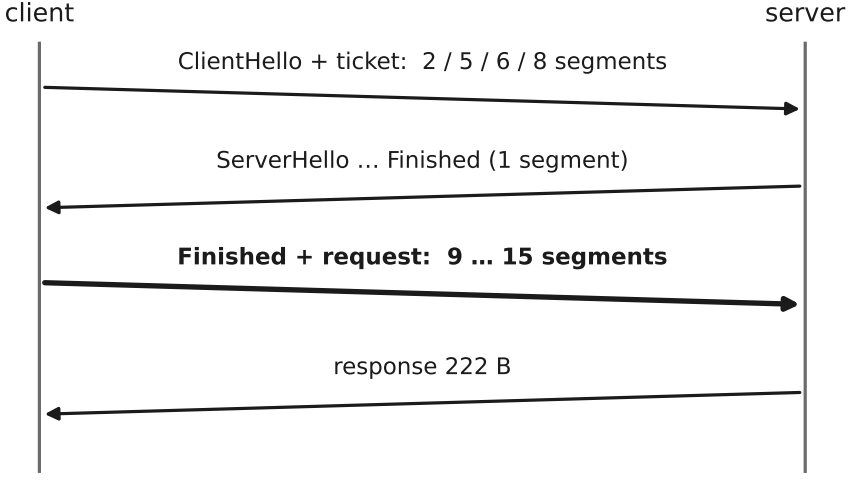

What is measured

resumed latency: ECDSA-P256, ML-DSA-44/-65/-87,

depth 2, RTT 100 ms, MTU 1500

ClientHello length = the client's window at the request

Compared with

slack_predict.py, slack = 0 … 3: written and committed before the run

28 cells; wrong = a whole round trip off (> 50 ms)